

regression-starterkit

July 3, 2018

1 Training regression model using gradient descent

1.1 Setup

```
In [1]: import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
from mpl_toolkits.mplot3d import Axes3D # required even though not ref'd!
import random
from matplotlib import rcParams

EXACT_COEFF = [-2977.3484892930505, 113.13288354356122]

rcParams["font.size"] = 12

def line_fit(B, x):
    return B[2-1]*x + B[1-1]
```

1.2 Load data


```
In [2]: df = pd.read_csv('data/cheese_deaths.csv')
df
```

```
Out[2]:
```


	years	cheese	deaths
0	2000	29.8	327
1	2001	30.1	456
2	2002	30.5	509
3	2003	30.6	497
4	2004	31.3	596
5	2005	31.7	573
6	2006	32.6	661
7	2007	32.7	809
8	2008	32.8	717
9	2009	33.1	741

1.3 Reproduce plots

```
In [4]: fig, ax1 = plt.subplots(figsize=(8,2.5))
...
plt.show()
```



```
In [3]: ...
```



1.4 Define cost function and gradient

```
In [5]: def MSE(B):
...     """
...     Compute and return the mean squared error with:
...     (x_i, y_i) = (df.cheese, df.deaths).
...     The line coefficients are B = [y-intercept, slope].
...     """
```

```
In [18]: def MSE_gradient(B):
...     """
```

Compute and return the gradient vector containing the partial derivatives of the MSE cost function with respect to b_1 and b_2 . Don't scale by 2 and divide by the number of observations as it's a waste of computation. We can adjust the learning rate to take this into consideration.

"""

1.5 Define minimize function

```
In [19]: def minimize(B0, eta, precision):
```

"""

Minimize the MSE cost function starting at model parameters B_0 and using learning rate vector η . The model has converged when the change in MSE from one iteration to the next is less than precision. Return both the B at which your algorithm converged and also the list of B points between B_0 and converged B , inclusively.

"""

```
    trace = []
```

```
In [8]: LEARNING_RATES = [0.02,0.0001]
```

```
    PRECISION = 0.000000001 # can't be too small as x-xprev prec is low
```

```
    B0 = np.array([random.randrange(-6000, 4000), random.randrange(-200, 300)])
```

```
In [20]: %time B,trace = minimize(B0, LEARNING_RATES, PRECISION)
```

```
CPU times: user 7.22 s, sys: 262 ms, total: 7.49 s
```

```
Wall time: 7.51 s
```

```
In [9]: print(f"descent gives {B} in {len(trace)} steps")
```

```
    print(f"exact is {EXACT_COEFF}")
```

```
    print(f"cost of approx is {MSE(B):1.10f}")
```

```
    print(f"cost of exact is {MSE(EXACT_COEFF):1.10f}")
```

```
descent gives [-2977.30834673  113.13161144] in 47004 steps
```

```
exact is      [-2977.3484892930505, 113.13288354356122]
```

```
cost of approx is 2004.4563237107
```

```
cost of exact  is 2004.4563215019
```

1.6 Draw heatmap with trace

```
In [10]: def get_surface(Cost, b1, b2):
```

"""

Given an arbitrary function of two scalar parameters, f , return a list of lists, C , containing f evaluated at $f(b1[i],b2[j])$. i is the b_1 coordinate and j is the b_2 coordinate. b_1 is the np.arange of y -intercept values and b_2 is the np.arange of slope values.

"""

```

In [13]: def heatmap(b0, b1, C, trace): # trace is a list of [b1, b2] pairs
        """
        Given b0 and b1, y-intercept and slope values, and cost 2D matrix, plot
        a heatmap using plt.imshow(). Make the heat map look like those in
        the project description.

        Once the heat map is displayed, plot the trace coordinates in color
        #FEE08F. Plot the first coordinate in color #FB060B and the last
        coordinate in black.

        Returns nothing.
        """
        fig = plt.figure(figsize=(6,6))
        ...
        plt.imshow(C,
            origin='lower',
            cmap='coolwarm',
            alpha=0.65,
            ...
        )
        ...
        # Plot all trace coordinates at once instead of using a Python loop
        # (MUCH faster this way)
        ax.plot(trace[:,0], trace[:,1], color='#FEE08F', marker='.', markersize=.1)
        ...

In [14]: b1 = np.arange(-6000, 4000, 30) # y intercept
        b2 = np.arange(-200, 300, 1)     # slope
        C = get_surface(MSE, b1, b2)
        heatmap(b1, b2, C, trace)

```

